


Chapter 2 Application Layer



*Computer Networking:
A Top Down Approach*,
6th edition,
Jim Kurose, Keith Ross
Addison-Wesley, Feb
2012.

2: Application Layer 1

Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - ✦ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with UDP
- 2.8 Socket programming with TCP

2: Application Layer 2

Chapter 2: Application Layer

Our goals:

- conceptual, implementation aspects of network application protocols
 - ✦ transport-layer service models
 - ✦ client-server paradigm
 - ✦ peer-to-peer paradigm
- learn about protocols by examining popular application-level protocols
 - ✦ HTTP
 - ✦ FTP
 - ✦ SMTP / POP3 / IMAP
 - ✦ DNS
- programming network applications
 - ✦ socket API

2: Application Layer 3

Some network apps

- e-mail
- web
- instant messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video clips
- social networks
- voice over IP
- real-time video conferencing
- grid computing

2: Application Layer 4


Creating a network app

write programs that

- ✦ run on (different) end systems
- ✦ communicate over network
- ✦ e.g., web server software communicates with browser software

No need to write software for network-core devices

- ✦ Network-core devices do not run user applications
- ✦ applications on end systems allows for rapid app development, propagation



2: Application Layer 5

Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - ✦ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with UDP
- 2.8 Socket programming with TCP


2: Application Layer 6

Application architectures

- Client-server
 - ❖ Including data centers/cloud computing
- Peer-to-peer (P2P)
- Hybrid of client-server and P2P

2: Application Layer 7

Client-server architecture



server:

- ❖ always-on host
- ❖ permanent IP address
- ❖ server farms for scaling

clients:

- ❖ communicate with server
- ❖ may be intermittently connected
- ❖ may have dynamic IP addresses
- ❖ do not communicate directly with each other

2: Application Layer 8

Google Data Centers


- Estimated cost of data center: \$600M
- Google spent \$2.4B in 2007 on new data centers



2: Application Layer 9

Pure P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses



Highly scalable but difficult to manage

2: Application Layer 10

Hybrid of client-server and P2P

Skype

- ❖ voice-over-IP P2P application
- ❖ centralized server: finding address of remote party:
- ❖ client-client connection: direct (not through server)

Instant messaging

- ❖ chatting between two users is P2P
- ❖ centralized service: client presence detection/location
 - user registers its IP address with central server when it comes online
 - user contacts central server to find IP addresses of buddies

2: Application Layer 11

App-layer protocol defines

- Types of messages exchanged,
 - ❖ e.g., request, response
- Message syntax:
 - ❖ what fields in messages & how fields are delineated
- Message semantics
 - ❖ meaning of information in fields
- Rules for when and how processes send & respond to messages

Public-domain protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP, BitTorrent

Proprietary protocols:

- e.g., Skype, ppsstream

2: Application Layer 12

What transport service does an app need?

Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer

Timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be "effective"

Throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be "effective"
- other apps ("elastic apps") make use of whatever throughput they get

Security

- Encryption, data integrity, ...

2: Application Layer 13

Transport service requirements of common apps

| Application | Data loss | Throughput | Time Sensitive |
|-----------------------|---------------|---|-----------------|
| file transfer | no loss | elastic | no |
| e-mail | no loss | elastic | no |
| Web documents | no loss | elastic | no |
| real-time audio/video | loss-tolerant | audio: 5kbps-1Mbps video: 10kbps-5Mbps | yes, 100's msec |
| stored audio/video | loss-tolerant | same as above | yes, few secs |
| interactive games | loss-tolerant | few kbps up | yes, 100's msec |
| instant messaging | no loss | elastic | yes and no |

2: Application Layer 14

Internet transport protocols services

TCP service:

- ❑ **connection-oriented**: setup required between client and server processes
- ❑ **reliable transport** between sending and receiving processes
- ❑ **flow control**: sender won't overwhelm receiver
- ❑ **congestion control**: throttle sender when network overloaded
- ❑ **does not provide**: timing, minimum throughput guarantees, security

UDP service:

- ❑ unreliable data transfer between sending and receiving process
- ❑ does not provide: connection setup, reliability, flow control, congestion control, timing, throughput guarantee, or security

Q: why bother? Why is there a UDP?

2: Application Layer 15

Internet apps: application, transport protocols

| | Application layer protocol | Underlying transport protocol |
|------------------------|--------------------------------------|-------------------------------|
| | e-mail | SMTP [RFC 2821] |
| remote terminal access | Telnet [RFC 854] | TCP |
| | Web | HTTP [RFC 2616] |
| | file transfer | FTP [RFC 959] |
| streaming multimedia | HTTP (eg Youtube), RTP [RFC 1889] | TCP or UDP |
| | Internet telephony (e.g., Skype) | typically UDP |

2: Application Layer 16

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 FTP
- ❑ 2.4 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- ❑ 2.5 DNS
- ❑ 2.6 P2P applications
- ❑ 2.7 Socket programming with UDP
- ❑ 2.8 Socket programming with TCP

2: Application Layer 17

Web and HTTP

First some jargon

- ❑ Web page consists of **objects**
- ❑ Object can be HTML file, JPEG image, Java applet, audio file,...
- ❑ Web page consists of **base HTML-file** which includes several referenced objects
- ❑ Each object is addressable by a **URL**
- ❑ Example URL:

www.someschool.edu/someDept/pic.gif


host name path name

2: Application Layer 18

HTTP overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - client**: browser that requests, receives, "displays" Web objects
 - server**: Web server sends objects in response to requests



2: Application Layer 19

HTTP overview (continued)

Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is "stateless"

- server maintains no information about past client requests

Protocols that maintain "state" are complex!

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled

2: Application Layer 20

HTTP connections

Nonpersistent HTTP

- At most one object is sent over a TCP connection.

Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.

2: Application Layer 21

Nonpersistent HTTP

Suppose user enters URL

`www.someSchool.edu/someDepartment/home.index` (contains text, references to 10 jpeg images)

- 1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80
- 1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. "accepts" connection, notifying client
2. HTTP client sends HTTP **request message** (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`
3. HTTP server receives request message, forms **response message** containing requested object, and sends message into its socket

time

2: Application Layer 22

Nonpersistent HTTP (cont.)

4. HTTP server closes TCP connection.
5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects

time


2: Application Layer 23

Non-Persistent HTTP: Response time

Definition of RTT: time for a small packet to travel from client to server and back.

Response time:

- one RTT to initiate TCP connection
 - one RTT for HTTP request and first few bytes of HTTP response to return
 - file transmission time
- total = 2RTT + transmit time**



2: Application Layer 24

Persistent HTTP

Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

2: Application Layer 25


HTTP request message

- two types of HTTP messages: **request**, **response**
- HTTP request message:**
 - ASCII (human-readable format)

request line (GET, POST, HEAD commands) → GET /somedir/page.html HTTP/1.1
 header lines → Host: www.someschool.edu
 User-agent: Mozilla/4.0
 Connection: close
 Accept-language: fr
 Carriage return line feed (extra carriage return, line feed) indicates end of message

2: Application Layer 26

HTTP request message: general format



2: Application Layer 27

Uploading form input

Post method:

- Web page often includes form input
- Input is uploaded to server in entity body

URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

www.somesite.com/animalsearch?monkeys&banana

2: Application Layer 28

Method types

HTTP/1.0

- GET
- POST
- HEAD
 - asks server to leave requested object out of response

HTTP/1.1

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field

2: Application Layer 29

HTTP response message

status line (protocol status code status phrase) → HTTP/1.1 200 OK
 header lines → Connection: close
 Date: Thu, 06 Aug 1998 12:00:15 GMT
 Server: Apache/1.3.0 (Unix)
 Last-Modified: Mon, 22 Jun 1998
 Content-Length: 6821
 Content-Type: text/html
 data, e.g., requested HTML file → data data data data data ...

2: Application Layer 30

HTTP response status codes

In first line in server->client response message.
A few sample codes:

- 200 OK**
 - ❖ request succeeded, requested object later in this message
- 301 Moved Permanently**
 - ❖ requested object moved, new location specified later in this message (Location:)
- 400 Bad Request**
 - ❖ request message not understood by server
- 404 Not Found**
 - ❖ requested document not found on this server
- 505 HTTP Version Not Supported**

2: Application Layer 31

User-server state: cookies

Many major Web sites use cookies

Four components:

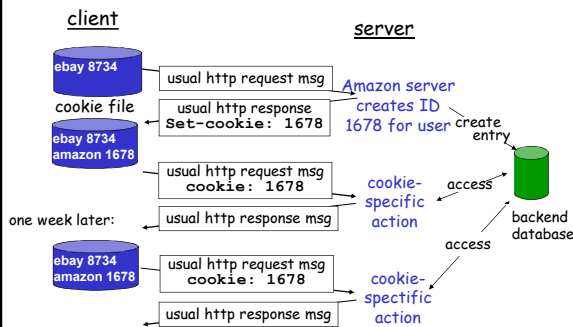
- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

- ❑ Susan always access Internet always from PC
- ❑ visits specific e-commerce site for first time
- ❑ when initial HTTP requests arrives at site, site creates:
 - ❖ unique ID
 - ❖ entry in backend database for ID

2: Application Layer 32

Cookies: keeping "state" (cont.)



2: Application Layer 33

Cookies (continued)

What cookies can bring:

- ❑ authorization
- ❑ shopping carts
- ❑ recommendations
- ❑ user session state (Web e-mail)

Cookies and privacy:

- ❑ cookies permit sites to learn a lot about you
- ❑ you may supply name and e-mail to sites

How to keep "state":


- ❑ protocol endpoints: maintain state at sender/receiver over multiple transactions
- ❑ cookies: http messages carry state

2: Application Layer 34

Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ❑ user sets browser: Web accesses via cache
- ❑ browser sends all HTTP requests to cache
 - ❖ object in cache: cache returns object
 - ❖ else cache requests object from origin server, then returns object to client



2: Application Layer 35

More about Web caching

- ❑ cache acts as both client and server
- ❑ typically cache is installed by ISP (university, company, residential ISP)

Why Web caching?

- ❑ reduce response time for client request
- ❑ reduce traffic on an institution's access link.
- ❑ Internet dense with caches: enables "poor" content providers to effectively deliver content (but so does P2P file sharing)

2: Application Layer 36


Caching example

Assumptions

- average object size = 1,000,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay = 2 sec + minutes + milliseconds



2: Application Layer 37


Caching example (cont)

possible solution

- increase bandwidth of access link to, say, 100 Mbps

consequence

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay = 2 sec + msec + msec
- often a costly upgrade



2: Application Layer 38


Caching example (cont)

possible solution: install cache

- suppose hit rate is 0.4

consequence

- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay = .6*(2.01) secs + .4*milliseconds < 1.4 secs




2: Application Layer 39

Conditional GET

- Goal:** don't send object if cache has up-to-date cached version
- cache:** specify date of cached copy in HTTP request
If-modified-since: <date>
- server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified

cache

server




2: Application Layer 40

Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with UDP
- 2.8 Socket programming with TCP

2: Application Layer 41

FTP: the file transfer protocol

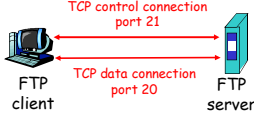


- transfer file to/from remote host
- client/server model
 - client: side that initiates transfer (either to/from remote)
 - server: remote host

- ftp: RFC 959
- ftp server: port 21

2: Application Layer 42

FTP: separate control, data connections

- FTP client contacts FTP server at port 21, TCP is transport protocol
 - client authorized over control connection
 - client browses remote directory by sending commands over control connection.
 - when server receives file transfer command, server opens 2nd TCP connection (for file) to client
 - after transferring one file, server closes data connection.
- 
- server opens another TCP data connection to transfer another file.
 - control connection: "out of band"
 - FTP server maintains "state": current directory, earlier authentication

2: Application Layer 43

FTP commands, responses

Sample commands:

- sent as ASCII text over control channel
- USER *username*
- PASS *password*
- LIST return list of file in current directory
- RETR *filename* retrieves (gets) file
- STOR *filename* stores (puts) file onto remote host

Sample return codes

- status code and phrase (as in HTTP)
- 331 Username OK, password required
- 125 data connection already open; transfer starting
- 425 Can't open data connection
- 452 Error writing file

2: Application Layer 44

Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with UDP
- 2.8 Socket programming with TCP

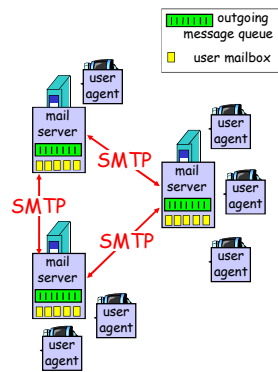
2: Application Layer 45

Electronic Mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

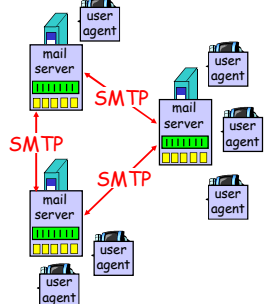
User Agent

- a.k.a. "mail reader"
 - composing, editing, reading mail messages
 - e.g., Eudora, Outlook, elm, Mozilla Thunderbird
 - outgoing, incoming messages stored on server
- 

2: Application Layer 46

Electronic Mail: mail servers

Mail Servers

- mailbox contains incoming messages for user
 - message queue of outgoing (to be sent) mail messages
 - SMTP protocol between mail servers to send email messages
 - client: sending mail server
 - "server": receiving mail server
- 

2: Application Layer 47


Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- direct transfer: sending server to receiving server
- three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- command/response interaction
 - commands: ASCII text
 - response: status code and phrase
- messages must be in 7-bit ASCII

2: Application Layer 48

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and "to" bob@someschool.edu
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



2: Application Layer 49

Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

2: Application Layer 50

SMTP: final words


- SMTP uses persistent connections
 - SMTP requires message (header & body) to be in 7-bit ASCII
 - SMTP server uses CRLF.CRLF to determine end of message
- Comparison with HTTP:**
- HTTP: pull
 - SMTP: push
 - both have ASCII command/response interaction, status codes
 - HTTP: each object encapsulated in its own response msg
 - SMTP: multiple objects sent in multipart msg

2: Application Layer 51

Mail message format


SMTP: protocol for exchanging email msgs
RFC 822: standard for text message format:

- header lines, e.g.,
 - To:
 - From:
 - Subject:
 different from SMTP commands!
- body
 - the "message", ASCII characters only



2: Application Layer 52

Mail access protocols



- SMTP: delivery/storage to receiver's server
- Mail access protocol: retrieval from server
 - POP: Post Office Protocol [RFC 1939]
 - authorization (agent <--> server) and download
 - IMAP: Internet Mail Access Protocol [RFC 1730]
 - more features (more complex)
 - manipulation of stored msgs on server
 - HTTP: gmail, Hotmail, Yahoo! Mail, etc.

2: Application Layer 53

POP3 protocol

authorization phase

- client commands:
 - user: declare username
 - pass: password
- server responses
 - +OK
 - ERR

transaction phase, client:

- list: list message numbers
- retr: retrieve message by number
- dele: delete
- quit

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

2: Application Layer 54

POP3 (more) and IMAP

More about POP3

- ❑ Previous example uses "download and delete" mode.
- ❑ Bob cannot re-read e-mail if he changes client
- ❑ "Download-and-keep": copies of messages on different clients
- ❑ POP3 is stateless across sessions

IMAP

- ❑ Keep all messages in one place: the server
- ❑ Allows user to organize messages in folders
- ❑ IMAP keeps user state across sessions:
 - ✦ names of folders and mappings between message IDs and folder name

2: Application Layer 55

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 FTP
- ❑ 2.4 Electronic Mail
 - ✦ SMTP, POP3, IMAP
- ❑ 2.5 DNS
- ❑ 2.6 P2P applications
- ❑ 2.7 Socket programming with UDP
- ❑ 2.8 Socket programming with TCP

2: Application Layer 56

DNS: Domain Name System

People: many identifiers:

- ✦ SSN, name, passport #

Internet hosts, routers:

- ✦ IP address (32 bit) - used for addressing datagrams
- ✦ "name", e.g., www.yahoo.com - used by humans

Q: map between IP addresses and name ?

Domain Name System:

- ❑ *distributed database* implemented in hierarchy of many *name servers*
- ❑ *application-layer protocol* host, routers, name servers to communicate to *resolve* names (address/name translation)
 - ✦ note: core Internet function, implemented as application-layer protocol
 - ✦ complexity at network's "edge"

2: Application Layer 57

DNS

DNS services

- ❑ hostname to IP address translation
- ❑ host aliasing
 - ✦ Canonical, alias names
- ❑ mail server aliasing
- ❑ load distribution
 - ✦ replicated Web servers: set of IP addresses for one canonical name


Why not centralize DNS?

- ❑ single point of failure
- ❑ traffic volume
- ❑ distant centralized database
- ❑ maintenance

doesn't scale!

2: Application Layer 58

Distributed, Hierarchical Database



Client wants IP for www.amazon.com; 1st approx:

- ❑ client queries a root server to find com DNS server
- ❑ client queries com DNS server to get amazon.com DNS server
- ❑ client queries amazon.com DNS server to get IP address for www.amazon.com

2: Application Layer 59

DNS: Root name servers

- ❑ contacted by local name server that can not resolve name
- ❑ root name server:
 - ✦ contacts authoritative name server if name mapping not known
 - ✦ gets mapping
 - ✦ returns mapping to local name server



2: Application Layer 60

TLD and Authoritative Servers

- **Top-level domain (TLD) servers:**
 - ❖ responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
 - ❖ Network Solutions maintains servers for com TLD
 - ❖ Educause for edu TLD
- **Authoritative DNS servers:**
 - ❖ organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
 - ❖ can be maintained by organization or service provider

2: Application Layer 61

Local Name Server

- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one.
 - ❖ also called "default name server"
- when host makes DNS query, query is sent to its local DNS server
 - ❖ acts as proxy, forwards query into hierarchy


2: Application Layer 62

DNS name resolution example

- Host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"




2: Application Layer 63

DNS name resolution example

recursive query:

- puts burden of name resolution on contacted name server
- heavy load?



2: Application Layer 64

DNS: caching and updating records

- once (any) name server learns mapping, it *caches* mapping
 - ❖ cache entries timeout (disappear) after some time
 - ❖ TLD servers typically cached in local name servers
 - Thus root name servers not often visited
- update/notify mechanisms under design by IETF
 - ❖ RFC 2136
 - ❖ <http://www.ietf.org/html.charters/dnsind-charter.html>

2: Application Layer 65

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

- | | |
|---|---|
| <ul style="list-style-type: none"> □ Type=A <ul style="list-style-type: none"> ❖ name is hostname ❖ value is IP address □ Type=NS <ul style="list-style-type: none"> ❖ name is domain (e.g. foo.com) ❖ value is hostname of authoritative name server for this domain | <ul style="list-style-type: none"> □ Type=CNAME <ul style="list-style-type: none"> ❖ name is alias name for some "canonical" (the real) name www.ibm.com is really servereast.backup2.ibm.com ❖ value is canonical name □ Type=MX <ul style="list-style-type: none"> ❖ value is name of mailserver associated with name |
|---|---|

2: Application Layer 66

DNS protocol, messages

DNS protocol: *query* and *reply* messages, both with same *message format*


msg header

- **identification**: 16 bit #
for query, reply to query uses same #
- **flags**:
 - ❖ query or reply
 - ❖ recursion desired
 - ❖ recursion available
 - ❖ reply is authoritative

| identification | flags |
|---|---|
| number of questions | number of answer RRs |
| number of authority RRs | number of additional RRs |
| questions (variable number of questions) | answers (variable number of resource records) |
| | authority (variable number of resource records) |
| | additional information (variable number of resource records) |

2: Application Layer 67

DNS protocol, messages



2: Application Layer 68

Inserting records into DNS

- example: new startup "Network Utopia"
- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
 - ❖ provide names, IP addresses of authoritative name server (primary and secondary)
 - ❖ registrar inserts two RRs into com TLD server:

```
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
```

- create authoritative server Type A record for www.networkutopia.com: Type MX record for networkutopia.com
- *How do people get IP address of your Web site?*

2: Application Layer 69


Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with UDP
- 2.8 Socket programming with TCP

2: Application Layer 70

Pure P2P architecture


- no always-on server
- arbitrary end systems directly communicate *peer-peer*
- peers are intermittently connected and change IP addresses
- **Three topics**:
 - ❖ File distribution
 - ❖ Searching for information
 - ❖ Case Study: Skype



2: Application Layer 71

File Distribution: Server-Client vs P2P


Question: How much time to distribute file from one server to N peers?



2: Application Layer 72

File distribution time: server-client

- server sequentially sends N copies:
 - NF/u_s time
- client i takes F/d_i time to download




Time to distribute F to N clients using client/server approach

$d_{cs} = \max \{ NF/u_s, F/\min(d_i) \}$
increases linearly in N (for large N)

2: Application Layer 73

File distribution time: P2P

- server must send one copy: F/u_s time
- client i takes F/d_i time to download
- NF bits must be downloaded (aggregate)
 - fastest possible upload rate: $u_s + \sum u_i$




$$d_{p2p} = \max \{ F/u_s, F/\min(d_i), NF/(u_s + \sum u_i) \}$$

2: Application Layer 74

Server-client vs. P2P: example

Client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{\min} \geq u_s$




2: Application Layer 75

File distribution: BitTorrent

- P2P file distribution

tracker: tracks peers participating in torrent


torrent: group of peers exchanging chunks of a file



2: Application Layer 76

BitTorrent (1)

- file divided into 256KB **chunks**.
- peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- while downloading, peer uploads chunks to other peers.
- peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain



2: Application Layer 77

BitTorrent (2)

Pulling Chunks

- at any given time, different peers have different subsets of file chunks
- periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- Alice sends requests for her missing chunks
 - rarest first


Sending Chunks: tit-for-tat

- Alice sends chunks to four neighbors currently sending her chunks at the **highest rate**
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - newly chosen peer may join top 4
 - "optimistically unchoke"

2: Application Layer 78

BitTorrent: Tit-for-tat

- (1) Alice "optimistically unchokes" Bob
- (2) Alice becomes one of Bob's top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice's top-four providers



2: Application Layer 79

Distributed Hash Table (DHT)

- DHT = distributed P2P database
- Database has (key, value) pairs;
 - ✦ key: ss number; value: human name
 - ✦ key: content type; value: IP address
- Peers **query** DB with key
 - ✦ DB returns values that match the key
- Peers can also **insert** (key, value) pairs

2: Application Layer 80

DHT Identifiers

- Assign integer identifier to each peer in range $[0, 2^n - 1]$.
 - ✦ Each identifier can be represented by n bits.
- Require each key to be an integer in **same range**.
- To get integer keys, hash original key.
 - ✦ eg, key = $h(\text{"Led Zeppelin IV"})$
 - ✦ This is why they call it a distributed "hash" table


2: Application Layer 81

How to assign keys to peers?

- Central issue:
 - ✦ Assigning (key, value) pairs to peers.
- Rule: assign key to the peer that has the **closest** ID.
- Convention in lecture: closest is the **immediate successor** of the key.
- Ex: $n=4$; peers: 1,3,4,5,8,10,12,14;
 - ✦ key = 13, then successor peer = 14
 - ✦ key = 15, then successor peer = 1

2: Application Layer 82

Circular DHT (1)




- Each peer **only** aware of immediate successor and predecessor.
- "Overlay network"

2: Application Layer 83

Circle DHT (2)


$O(N)$ messages on avg to resolve query, when there are N peers



Define **closest** as closest successor

2: Application Layer 84


Circular DHT with Shortcuts



- Each peer keeps track of IP addresses of predecessor, successor, short cuts.
- Reduced from 6 to 2 messages.
- Possible to design shortcuts so $O(\log N)$ neighbors, $O(\log N)$ messages in query

2: Application Layer 85

Peer Churn




- To handle peer churn, require each peer to know the IP address of its two successors.
- Each peer periodically pings its two successors to see if they are still alive.

- Peer 5 abruptly leaves
- Peer 4 detects; makes 8 its immediate successor; asks 8 who its immediate successor is; makes 8's immediate successor its second successor.
- What if peer 13 wants to join?

2: Application Layer 86

P2P Case study: Skype


- inherently P2P: pairs of users communicate.
- proprietary application-layer protocol (inferred via reverse engineering)
- hierarchical overlay with SNs
- Index maps usernames to IP addresses; distributed over SNs



2: Application Layer 87

Peers as relays

- Problem when both Alice and Bob are behind "NATs".
 - NAT prevents an outside peer from initiating a call to insider peer
- Solution:
 - Using Alice's and Bob's SNs, Relay is chosen
 - Each peer initiates session with relay.
 - Peers can now communicate through NATs via relay



2: Application Layer 88

Chapter 2: Summary

our study of network apps now complete!

- application architectures
 - client-server
 - P2P
 - hybrid
- application service requirements:
 - reliability, bandwidth, delay
- Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- specific protocols:
 - HTTP
 - FTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent, Skype
- socket programming

2: Application Layer 89

Chapter 2: Summary

Most importantly: learned about protocols

- typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- message formats:
 - headers: fields giving info about data
 - data: info being communicated
- Important themes:
 - control vs. data msgs
 - in-band, out-of-band
 - centralized vs. decentralized
 - stateless vs. stateful
 - reliable vs. unreliable msg transfer
 - "complexity at network edge"

2: Application Layer 90