

Alpha-Beta Pruning: Algorithm and Analysis

Tsan-sheng Hsu

徐讚昇

tshsu@iis.sinica.edu.tw

<http://www.iis.sinica.edu.tw/~tshsu>

1


Introduction

- Alpha-beta pruning is the standard searching procedure used for solving 2-person perfect-information zero sum games exactly.
- Definitions:
 - A position p .
 - The **value** of a position p , $f(p)$, is a numerical value computed from evaluating p .
 - Value is computed from the root player's point of view.
 - Positive values mean in favor of the root player.
 - Negative values mean in favor of the opponent.
 - Since it is a zero sum game, thus from the opponent's point of view, the value can be assigned $-f(p)$.
 - A **terminal position**: a position whose value can be decided.
 - A position where win/loss/draw can be concluded.
 - A position where some constraints are met.
 - A position p has b legal moves p_1, p_2, \dots, p_b .

TCG: α - β Pruning, 20181102, Tsan-sheng Hsu ©

2

Tree node numbering




- From the root, number a node in a search tree by a sequence of integers $a_1, a_2, a_3, a_4, \dots$
 - Meaning from the root, you first take the a_1 th branch, then the a_2 th branch, and then the a_3 th branch, and then the a_4 th branch \dots
 - The root is specified as an empty sequence.
 - The **depth** of a node is the length of the sequence of integers specifying it.
- This is called "Dewey decimal system."

TCG: α - β Pruning, 20181102, Tsan-sheng Hsu ©

3

Mini-max formulation




Mini-max formulation:

- $$F'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \max\{G'(p_1), \dots, G'(p_b)\} & \text{if } b > 0 \end{cases}$$
- $$G'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \min\{F'(p_1), \dots, F'(p_b)\} & \text{if } b > 0 \end{cases}$$
- An indirect recursive formula with a bottom-up evaluation!
- Equivalent to AND-OR logic.

TCG: α - β Pruning, 20181102, Tsan-sheng Hsu ©

4

Mini-max formulation




Mini-max formulation:

- $$F'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \max\{G'(p_1), \dots, G'(p_b)\} & \text{if } b > 0 \end{cases}$$
- $$G'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \min\{F'(p_1), \dots, F'(p_b)\} & \text{if } b > 0 \end{cases}$$
- An indirect recursive formula with a bottom-up evaluation!
- Equivalent to AND-OR logic.

TCG: α - β Pruning, 20181102, Tsan-sheng Hsu ©

5

Mini-max formulation




Mini-max formulation:

- $$F'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \max\{G'(p_1), \dots, G'(p_b)\} & \text{if } b > 0 \end{cases}$$
- $$G'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \min\{F'(p_1), \dots, F'(p_b)\} & \text{if } b > 0 \end{cases}$$
- An indirect recursive formula with a bottom-up evaluation!
- Equivalent to AND-OR logic.

TCG: α - β Pruning, 20181102, Tsan-sheng Hsu ©

6

Mini-max formulation



Mini-max formulation:

- $$F'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \max\{G'(p_1), \dots, G'(p_b)\} & \text{if } b > 0 \end{cases}$$
- $$G'(p) = \begin{cases} f(p) & \text{if } b = 0 \\ \min\{F'(p_1), \dots, F'(p_b)\} & \text{if } b > 0 \end{cases}$$
- An indirect recursive formula with a bottom-up evaluation!
- Equivalent to AND-OR logic.

Mini-max: revised (1/2)

- Algorithm $F'(\text{position } p, \text{integer depth})$ // max node
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$ // a terminal node
 - or $\text{depth} = 0$ // remaining depth to search
 - or time is running up // from timing control
 - or some other constraints are met // add knowledge here
 - then return $f(p)$ // current board value
 - else begin
 - $m := -\infty$ // initial value
 - for $i := 1$ to b do // try each child
 - begin
 - $t := G'(p_i, \text{depth} - 1)$
 - if $t > m$ then $m := t$ // find max value
 - end
 - end
 - return m

Algorithm: Mini-max

- Algorithm $F'(\text{position } p)$ // max node
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - $m := -\infty$
 - for $i := 1$ to b do
 - $t := G'(p_i)$
 - if $t > m$ then $m := t$ // find max value
 - end;
 - return m
- Algorithm $G'(\text{position } p)$ // min node
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - $m := \infty$
 - for $i := 1$ to b do
 - $t := F'(p_i)$
 - if $t < m$ then $m := t$ // find min value
 - end;
 - return m


Mini-max: revised (2/2)

- Algorithm $G'(\text{position } p, \text{integer depth})$ // min node
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$ // a terminal node
 - or $\text{depth} = 0$ // remaining depth to search
 - or time is running up // from timing control
 - or some other constraints are met // add knowledge here
 - then return $f(p)$ // current board value
 - else begin
 - $m := \infty$ // initial value
 - for $i := 1$ to b do // try each child
 - begin
 - $t := F'(p_i, \text{depth} - 1)$
 - if $t < m$ then $m := t$ // find min value
 - end
 - end
 - return m

Mini-max: comments

- A **brute-force** method to try all possibilities!
 - May visit a position many times.
- Depth-first search
 - Move ordering is according to order the successor positions are generated.
 - Bottom-up evaluation.
 - Post-ordering traversal.
- Q:
 - Iterative deepening?
 - BFS?
 - Other types of searching?

Nega-max formulation




- Nega-max formulation:
 - Let $F(p)$ be the greatest possible value achievable from the position p against the optimal defensive strategy.

$$F(p) = \begin{cases} h(p) & \text{if } b = 0 \\ \max\{-F(p_1), \dots, -F(p_b)\} & \text{if } b > 0 \end{cases}$$

>

$$h(p) = \begin{cases} f(p) & \text{if depth of } p \text{ is 0 or even} \\ -f(p) & \text{if depth of } p \text{ is odd} \end{cases}$$

Nega-max formulation



- **Nega-max formulation:**
Let $F(p)$ be the greatest possible value achievable from position p against the optimal defensive strategy.

$$F(p) = \begin{cases} h(p) & \text{if } b = 0 \\ \max\{-F(p_1), \dots, -F(p_b)\} & \text{if } b > 0 \end{cases}$$


▷

$$h(p) = \begin{cases} f(p) & \text{if depth of } p \text{ is 0 or even} \\ -f(p) & \text{if depth of } p \text{ is odd} \end{cases}$$

Algorithm: Nega-max

- **Algorithm $F(\text{position } p, \text{ integer } \text{depth})$**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$ // a terminal node
 - or $\text{depth} = 0$ // remaining depth to search
 - or time is running up // from timing control
 - or some other constraints are met // add knowledge here
 - then return $h(p)$ else
 - begin
 - ▷ $m := -\infty$
 - ▷ for $i := 1$ to b do
 - ▷ begin
 - ▷ $t := -F(p_i, \text{depth}-1)$ // recursive call, the returned value is negated
 - ▷ if $t > m$ then $m := t$ // always find a max value
 - ▷ end
 - end
 - return m

Nega-max formulation



- **Nega-max formulation:**
Let $F(p)$ be the greatest possible value achievable from position p against the optimal defensive strategy.

$$F(p) = \begin{cases} h(p) & \text{if } b = 0 \\ \max\{-F(p_1), \dots, -F(p_b)\} & \text{if } b > 0 \end{cases}$$


▷

$$h(p) = \begin{cases} f(p) & \text{if depth of } p \text{ is 0 or even} \\ -f(p) & \text{if depth of } p \text{ is odd} \end{cases}$$

Nega-max: comments

- **Another brute-force method to try all possibilities.**
 - Use $h(p)$ instead of $f(p)$.
 - Watch out the code in dealing with search termination conditions.
 - ▷ Reach a given searching depth.
 - ▷ Timing control.
 - ▷ Other constraints such as the score is good or bad enough.
- **Notations:**
 - F' means the Mini-max version.
 - ▷ Need a G' companion.
 - ▷ Easy to explain.
 - F means the Negamax version.
 - ▷ Simpler code.
 - ▷ Maybe difficult to explain.

Nega-max formulation



- **Nega-max formulation:**
Let $F(p)$ be the greatest possible value achievable from position p against the optimal defensive strategy.

$$F(p) = \begin{cases} h(p) & \text{if } b = 0 \\ \max\{-F(p_1), \dots, -F(p_b)\} & \text{if } b > 0 \end{cases}$$


▷

$$h(p) = \begin{cases} f(p) & \text{if depth of } p \text{ is 0 or even} \\ -f(p) & \text{if depth of } p \text{ is odd} \end{cases}$$

Intuition for improvements


- **Branch-and-bound:** using information you have so far to **cut** or **prune** branches.
 - A branch is cut means we do not need to search it anymore.
 - If you know for sure or almost sure the value of your result is more than x and the current search result for this branch so far can give you no more than x ,
 - ▷ then there is no need to search this branch any further.
- **Two types of approaches**
 - Exact algorithms: through mathematical proof, it is guaranteed that the branches pruned **won't** contain the solution.
 - ▷ Alpha-beta pruning: reinvented by several researchers in the 1950's and 1960's.
 - ▷ Scout.
 - ▷ ...
 - Approximated heuristics: with a high probability that the solution won't be contained in the branches pruned.
 - ▷ Obtain a good estimation on the remaining cost.
 - ▷ Cut a branch when it is in a very bad position and there is little hope to gain back the advantage.

Alpha cut-off




- On the **max** node which is the root:
 - Assume you have finished exploring the branch at 1 and obtained the best value from it as *bound*.
 - You now search the branch at 2 by first searching the branch at 2.1.
 - Assume branch at 2.1 returns a value that is \leq *bound*.
 - Then no need to evaluate the branch at 2.2 and all later branches of 2, if any, at all.
 - The best possible value for the branch at 2 must be \leq *bound*.
 - Hence we should take value returned from the branch at 1 as the best possible solution.

Illustration — Deep alpha cut-off



Beta cut-off



- On the **min** node 1:
 - Assume you have finished exploring the branch at 1.1 and obtained the best value from it as *bound*.
 - You now search the branch at 1.2 by first exploring the branch at 1.2.1.
 - Assume the branch at 1.2.1 returns a value that is \geq *bound*.
 - Then no need to evaluate the branch at 1.2.2 and all later branches of 1.2, if any, at all.
 - The best possible value for the branch at 1.2 is \geq *bound*.
 - Hence we should take value returned from the branch at 1.1 as the best possible solution.

Ideas for refinements

- During searching, maintain two values *alpha* and *beta* so that
 - alpha* is the current lower bound of the possible returned value;
 - This means to say you know a way to achieve the value *alpha*.
 - beta* is the current upper bound of the possible returned value.
 - This means to say your opponent knows a way to achieve a value of *beta*.
- If during searching, we know for sure $\alpha > \beta$, then there is no need to search any more in this branch.
 - The returned value cannot be in this branch.
 - Backtrack until it is the case $\alpha \leq \beta$.
- The two values *alpha* and *beta* are called the ranges of the **current search window**.
 - These values are dynamic.
 - Initially, *alpha* is $-\infty$ and *beta* is ∞ .

Deep alpha cut-off

- For alpha cut-off:**
 - For a min node *u*, a branch of its ancestor (e.g., an elder brother of its parent) produces a lower bound V_i .
 - The first branch of *u* produces an upper bound V_u for *v*.
 - If $V_i \geq V_u$, then there is no need to evaluate the second branch and all later branches, of *u*.
- Deep alpha cut-off:**
 - Def: For a node *u* in a tree and a positive integer *g*, *Ancestor*(*g*, *u*) is the direct ancestor of *u* by tracing the parent's link *g* times.
 - When the lower bound V_i is produced at and propagated from *u*'s great grand parent, i.e., *Ancestor*(3,*u*), or any *Ancestor*(2*i*+1,*u*), *i* \geq 1.
 - When an upper bound V_u is returned from the a branch of *u* and $V_i \geq V_u$, then there is no need to evaluate all later branches of *u*.
- We can find similar properties for deep beta cut-off.


Alpha-beta pruning algorithm: Mini-Max

- Algorithm $F1'$ (position *p*, value *alpha*, value *beta*) // max node**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - $m := \alpha$
 - for $i := 1$ to b do
 - $t := G1'(p_i, m, \beta)$
 - if $t > m$ then $m := t$ // improve the current best value
 - if $m \geq \beta$ then return(β) // beta cut off
 - end; return *m*
- Algorithm $G1'$ (position *p*, value *alpha*, value *beta*) // min node**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - $m := \beta$
 - for $i := 1$ to b do
 - $t := F1'(p_i, \alpha, m)$
 - if $t < m$ then $m := t$
 - if $m \leq \alpha$ then return(α) // alpha cut off
 - end; return *m*

Example

Initial call: $F1'(\text{root}, -\infty, \infty)$


- $m = -\infty$
- call $G1'(\text{node } 1, -\infty, \infty)$
 - ▷ it is a terminal node
 - ▷ return value 15
- $t = 15$;
 - ▷ since $t > m$, m is now 15
- call $G1'(\text{node } 2, 15, \infty)$
 - ▷ call $F1'(\text{node } 2, 15, \infty)$
 - ▷ it is a terminal node; return 10
 - ▷ $t = 10$; since $t < \infty$, m is now 10
 - ▷ alpha is 15, m is 10, so we have an alpha cut off.
 - ▷ no need to call $F1'(\text{node } 2, 2, 15, 10)$
 - ▷ return 15
 - ▷ ...




Alpha-beta pruning algorithm: Nega-max

- **Algorithm $F1(\text{position } p, \text{value } \alpha, \text{value } \beta, \text{integer } depth)$**
 - **select the successor positions p_1, \dots, p_b**
 - **if $b = 0$ // a terminal node**
 - **$depth = 0$ // remaining depth to search**
 - **or time is running up // from timing control**
 - **or some other constraints are met // add knowledge here**
 - **then return $h(p)$ else**
 - **begin**
 - ▷ **$m := \alpha$**
 - ▷ **for $i := 1$ to b do**
 - ▷ **begin**
 - ▷ **$t := -F1(p_i, -\beta, -\alpha, depth - 1)$**
 - ▷ **if $t > m$ then $m := t$**
 - ▷ **if $m \geq \beta$ then return(β) // cut off**
 - ▷ **end**
 - **end**
 - **return m**


A complete example




Examples (1/4)




A complete example



Examples (2/4)




Examples (3/4)



Lessons from the previous examples

- It looks like for the same tree, **different move orderings** give very different cut branches.
- It looks like **if a node can evaluate a child with the best possible outcome earlier**, then it has a chance to cut earlier.
 - For a min node, this means **to search the child branch that gives the lowest value first**.
 - For a max node, this means **to search the child branch that gives the highest value first**.
- Comments:
 - It is impossible to always know which best branch is, or we do not have to do a brute-force search.
 - Q: In the best case scenario, how many nodes can be cut?


Examples (3/4)



Analysis of a possible best case

- Definitions:
 - A path in a search tree is a sequence of numbers indicating the branches selected in each level using the Dewey decimal system.
 - A position is denoted as a path $a_1.a_2.\dots.a_\ell$ from the root.
 - A position $a_1.a_2.\dots.a_\ell$ is **critical** if
 - ▷ $a_i = 1$ for all even values of i or
 - ▷ $a_i = 1$ for all odd values of i .
 - Note: as a special case, the root is critical.
 - Examples:
 - ▷ 2.1.4.1.2, 1.3.1.5.1.2, 1.1.1.2.1.1.1.3 and 1.1 are critical
 - ▷ 1.2.1.1.2 is not critical
- Q: Why the root needs to be critical?

Examples (4/4)



Perfect-ordering tree

- A **perfect-ordering tree**:


$$F(a_1, \dots, a_\ell) = \begin{cases} h(a_1, \dots, a_\ell) & \text{if } a_1, \dots, a_\ell \text{ is a terminal} \\ -F(a_1, \dots, a_\ell, 1) & \text{otherwise} \end{cases}$$
 - The first successor of every non-terminal position gives the best possible value.

Theorem 1

- **Theorem 1:** *F1* examines precisely the critical positions of a perfect-ordering tree.
- **Proof sketch:**
 - Classify the critical positions, a.k.a. nodes, into different types.
 - ▷ You must evaluate the first branch from the root to the bottom.
 - ▷ Alpha cut off happens at odd-depth nodes as soon as the first branch of this node is evaluated.
 - ▷ Beta cut off happens at even-depth nodes as soon as the first branch of this node is evaluated.
 - For nodes of the same type, associate them with pruning of same characteristics occurred.

Type 1 nodes

- **type 1: the root, or a node with all the a_i are 1;**
 - This means j does not exist.
 - Nodes on the leftmost branch.
 - **The leftmost child of a type 1 node except the root.**




Types of nodes

- **Classification of critical positions** $a_1.a_2.\dots.a_j.\dots.a_\ell$ where j is the least index, if exists, such that $a_j \neq 1$ and ℓ is the last index.
 - j will be the **anchor** in the analysis.
 - Def: let $IS1(a_i)$ be a boolean function so that it is 0 if it is not the value 1 and it is 1 if it is.
 - ▷ We call this **IS1 parity of a number**.
 - If j exists and $\ell > j$, then
 - ▷ $a_{j+1} = 1$ because this **position is critical and thus the IS1 parities of a_j and a_{j+1} are different**.
 - Since this position is critical, if $a_j \neq 1$, then $a_h = 1$ for any h such that $h - j$ is odd.
- We now classify critical nodes into 3 types.
 - Nodes of the same type share some common properties.

Type 2 nodes

- **Classification of critical positions** $a_1, a_2, \dots, a_j, \dots, a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- **type 2: $\ell - j$ is zero or even;**
 - **type 2.1: $\ell - j = 0$ which means $\ell = j$.**
 - ▷ It is in the form of $\underline{1, 1, 1, \dots, 1, 1, 1}, a_\ell$ and $a_\ell \neq 1$.
 - ▷ The non-leftmost children of a type 1 node.
 - **type 2.2: $\ell - j > 0$ and is even.**
 - ▷ It is in the form of $\underline{1, 1, \dots, 1, 1, a_j, 1, a_{j+2}, \dots, a_{\ell-2, 1}, a_\ell}$.
 - ▷ Note, we have already defined $\underline{1, 1, \dots, 1, 1, a_j, 1, a_{j+2}, \dots, a_{\ell-2, 1}}$ to be a type 3 node.
 - ▷ All of the children of a type 3 node.

Illustration — critical nodes



Type 3 nodes

- **Classification of critical positions** $a_1, a_2, \dots, a_j, \dots, a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- **type 3:** $\ell - j$ is odd;
 - $a_j \neq 1$ and $\ell - j$ is odd
 - ▷ Since this position is critical, the $IS1$ parities of a_j and a_ℓ are different.
 - $\implies a_\ell = 1$
 - $\implies a_{j+1} = 1$
 - It is in the form of
 - ▷ $1.1. \dots .1.a_j.1.a_{j+2}.1. \dots .1.a_{\ell-1}.1$.
 - The leftmost child of a **type 2 node**.
- **type 3.1:** $\ell - j = 1$.
 - ▷ It is of the form $1.1. \dots .1.a_j.1$
 - ▷ The leftmost child of a type 2.1 node.
- **type 3.2:** $\ell - j > 1$.
 - ▷ It is of the form $1.1. \dots .1.a_j.1.a_{j+2}.1. \dots .1.a_{\ell-1}.1$
 - ▷ The leftmost child of a type 2.2 node.

Comments

- Nodes of the same type have common properties.
- These properties can be used in solving other problems.
 - Example: Efficient parallelization of alpha-beta based searching algorithms.
- Main techniques used:
 - You cannot have two consecutive non-1 numbers in the ID of a critical node.
 - For each non-1 number, any number appeared later and is odd distance away must be 1.

Type 2.2 nodes

- Classification of critical positions $a_1.a_2.\dots.a_j.\dots.a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- type 2: $\ell - j$ is zero or even;
 - type 2.2: $\ell - j > 0$ and is even.
 - The IS1 parties of a_j and a_{j+1} are different.
 - \Rightarrow Since $a_j \neq 1$, $a_{j+1} = 1$.
 - $(\ell - 1) - j$ is odd:
 - \Rightarrow The IS1 parties of $a_{\ell-1}$ and a_j are different.
 - \Rightarrow Since $a_j \neq 1$, $a_{\ell-1} = 1$.
 - It is in the form of $1.1.\dots.1.1.a_j.1.a_{j+2}.\dots.a_{\ell-2}.1.a_\ell$.
 - Note, we will show $1.1.\dots.1.1.a_j.1.a_{j+2}.\dots.a_{\ell-2}.1$ is a type 3 node later.
 - All of the children of a type 3 node.

Type 2.1 nodes

- Classification of critical positions $a_1.a_2.\dots.a_j.\dots.a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- type 2: $\ell - j$ is zero or even;
 - type 2.1: $\ell - j = 0$.
 - Then $\ell = j$.
 - It is in the form of $1.1.1.\dots.1.1.1.a_\ell$ and $a_\ell \neq 1$.
 - The non-leftmost children of a type 1 node.






Illustration: Type 2.2 nodes



Type 3.1 nodes

- Classification of critical positions $a_1.a_2.\dots.a_j.\dots.a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- type 3: $\ell - j$ is odd;
 - type 3.1: $\ell - j = 1$.
 - It is of the form $1.1.\dots.1.a_j.1$ and $a_\ell \neq 1$.
 - The leftmost child of a type 2.1 node.



Type 3.2 nodes

- Classification of critical positions $a_1.a_2.\dots.a_j.\dots.a_\ell$ where j is the least index such that $a_j \neq 1$ and ℓ is the last index.
- type 3: $\ell - j$ is odd;
 - type 3.2: $\ell - j > 1$.
 - It is of the form $1.1.\dots.1.a_j.1.a_{j+2}.1.\dots.1.a_{\ell-1}.1$.
 - The leftmost child of a type 2.2 node.

Illustration: Type 3.2 nodes




Illustration of all nodes




Illustration of all nodes




Illustration of all nodes




Illustration of all nodes




Illustration of all nodes





Illustration of all nodes



Analysis: average case

- Assumptions: Let a random game tree be generated in such a way that each position on level j
 - has a probability q_j of being nonterminal and
 - has an average of b_j successors.

- Properties of the above random game tree
 - Expected number of positions on level ℓ is $b_0 \cdot b_1 \cdots b_{\ell-1}$
 - Expected number of positions on level ℓ examined by an alpha-beta procedure assumed the random game tree is perfectly ordered is

$$b_0 q_1 b_2 q_3 \cdots b_{\ell-2} q_{\ell-1} + q_0 b_1 q_2 b_3 \cdots q_{\ell-2} b_{\ell-1} - q_0 q_1 \cdots q_{\ell-1} \text{ if } \ell \text{ is even;}$$

$$b_0 q_1 b_2 q_3 \cdots q_{\ell-2} b_{\ell-1} + q_0 b_1 q_2 b_3 \cdots b_{\ell-2} q_{\ell-1} - q_0 q_1 \cdots q_{\ell-1} \text{ if } \ell \text{ is odd}$$


- Proof sketch:
 - If x is the expected number of positions of a certain type on level j , then $x \times b_j$ is the expected number of successors of these positions, and $x \times q_j$ is the expected number of "numbered 1" successors.
 - The above numbers equal to those of Corollary 1 when $q_j = 1$ and $b_j = b$ for $0 \leq j < \ell$.

Theorem 1: Proof sketch

- Properties (invariants)
 - A type 1 position p is examined by calling $F1(p, -\infty, \infty, \text{depth})$
 - p 's first successor p_1 is of type 1
 - $F(p) = -F(p_1) \neq \pm\infty$
 - p 's other successors p_2, \dots, p_b are of type 2
 - $p_i, i > 1$, are examined by calling $F1(p_i, -\infty, F(p_1), \text{depth})$
 - A type 2 position p is examined by calling $F1(p, -\infty, \text{beta}, \text{depth})$ where $-\infty < \text{beta} \leq F(p)$
 - p 's first successor p_1 is of type 3
 - $F(p) = -F(p_1)$
 - p 's other successors p_2, \dots, p_b are not examined
 - A type 3 position p is examined by calling $F1(p, \text{alpha}, \infty, \text{depth})$ where $\infty > \text{alpha} \geq F(p)$
 - p 's successors p_1, \dots, p_b are of type 2
 - they are examined by calling $F1(p_1, -\infty, -\text{alpha}, \text{depth}), \dots, F1(p_b, -\infty, -\text{max}\{m_{i-1}, \text{alpha}\}, \text{depth})$ where $m_i = F1(p_i, -\infty, -\text{max}\{m_{i-1}, \text{alpha}\}, \text{depth})$
- Using an inductive argument to prove.

Perfect ordering is not always the best

- Intuitively, we may "think" alpha-beta pruning would be most effective when a game tree is perfectly ordered.
 - That is, when the first successor of every position is the best possible move.
 - This is not always the case!



- Truly optimum order of game trees traversal is not obvious.

Analysis: best case

- Corollary 1: Assume each position has exactly b successors
 - The number of positions examined by the alpha-beta procedure on level i is exactly $b^{\lceil i/2 \rceil} + b^{\lfloor i/2 \rfloor} - 1$.
- Proof:
 - There are $b^{\lfloor i/2 \rfloor}$ sequences of the form a_1, \dots, a_i with $1 \leq a_i \leq b$ for all i such that $a_i = 1$ for all odd values of i .
 - There are $b^{\lceil i/2 \rceil}$ sequences of the form a_1, \dots, a_i with $1 \leq a_i \leq b$ for all i such that $a_i = 1$ for all even values of i .
 - We subtract 1 for the sequence 1.1.1.1.1 which are counted twice.
- Total number of nodes visited is

$$\sum_{i=0}^{\ell} b^{\lceil i/2 \rceil} + b^{\lfloor i/2 \rfloor} - 1.$$

When is a branch pruned?

- Assume a node r has two children u and v with u being visited before v using some move ordering.
 - Further assume u produced a new bound bound .
- Assume node v has a child w .
 - If the value new returned from w can cause a range conflict with bound , then branches of v later than w are cut.
- This means as long as the "relative" ordering of u and v are good enough, then we can have some cut-off.
 - There is no need for r to have the best move ordering.

Theorem 2

- **Theorem 2:** Alpha-beta pruning is optimum in the following sense:
 - Given any game tree and any algorithm which computes the value of the root position, there is a way to permute the tree
 - ▷ by reordering successor positions if necessary;
 - so that every terminal position examined by the alpha-beta method under this permutation is examined by the given algorithm.
 - Furthermore if the value of the root is not ∞ or $-\infty$, the alpha-beta procedure examines precisely the positions which are critical under this permutation.


Properties and comments

- **Properties:**
 - **Assumptions:** (1) $\alpha < \beta$ and (2) p is not a leaf.
 - $F1(p, \alpha, \beta, \text{depth}) = \alpha$ if $F(p) \leq \alpha$
 - $F1(p, \alpha, \beta, \text{depth}) = F(p)$ if $\alpha < F(p) < \beta$
 - $F1(p, \alpha, \beta, \text{depth}) = \beta$ if $F(p) \geq \beta$
 - $F1(p, -\infty, +\infty, \text{depth}) = F(p)$
- **Comments:**
 - $F1(p, \alpha, \beta, \text{depth})$: find the best possible value according to a **nega-max formula** for the position p with the constraints that
 - ▷ If $F(p) \leq \alpha$, then $F1(p, \alpha, \beta, \text{depth})$ returns with the value α from a terminal position whose value is $\leq \alpha$.
 - ▷ If $F(p) \geq \beta$, then $F1(p, \alpha, \beta, \text{depth})$ returns the value β from a terminal position whose value is $\geq \beta$.
 - **The meanings of α and β during searching:**
 - ▷ For a max node: the current best value is at least α .
 - ▷ For a min node: the current best value is at most β .
 - $F1$ always finds a value that is within α and β .
 - ▷ The bounds are hard, i.e., cannot be violated.

Variations of alpha-beta search

- Initially, to search a tree with the root r by calling $F1(r, -\infty, +\infty, \text{depth})$.
 - What does it mean to search a tree with the root r by calling $F1(r, \alpha, \beta, \text{depth})$?
 - ▷ To search the tree rooted at r requiring that the returned value to be within α and β .
- In an alpha-beta search with a pre-assigned window $[\alpha, \beta]$:
 - **Failed-high** means it returns a value that is larger than or equal to its upper bound β .
 - **Failed-low** means it returns a value that is smaller than or equal to its lower bound α .
- **Variations:**
 - **Brute force Nega-Max version:** F
 - ▷ Always finds the correct answer according to the Nega-Max formula.
 - **Original alpha-beta cut (Nega-Max) version:** $F1$
 - **Fail hard alpha-beta cut (Nega-Max) version:** $F2$
 - **Fail soft alpha-beta cut (Nega-Max) version:** $F3$

Original version: Example



- As long as the value of the leaf node W is less than the current α value, the returned value of A will be α .
- If the value of the leaf node W is greater than the current β value, the returned value of A will be β .

Original version

- Requiring $\alpha < \beta$
- **Algorithm $F1(\text{position } p, \text{value } \alpha, \text{value } \beta, \text{integer } \text{depth})$**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$ // a terminal node
 - or $\text{depth} = 0$ // remaining depth to search
 - or time is running up // from timing control
 - or some other constraints are met // add knowledge here
 - then return $h(p)$ else
 - begin
 - ▷ $m := \alpha$ // hard initial value
 - ▷ for $i := 1$ to b do
 - ▷ begin
 - ▷ $t := -F1(p_i, -\beta, -m, \text{depth} - 1)$
 - ▷ if $t > m$ then $m := t$ // the returned value is "used"
 - ▷ if $m \geq \beta$ then return(β) // cut off and return the hard bound
 - ▷ end
 - end
 - return m // if nothing over α , then α is returned

Alpha-beta pruning algorithm: Fail hard

- **Algorithm $F2'(\text{position } p, \text{value } \alpha, \text{value } \beta)$ // max node**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - ▷ $m := \alpha$
 - ▷ for $i := 1$ to b do
 - ▷ $t := G2'(p_i, m, \beta)$
 - ▷ if $t > m$ then $m := t$
 - ▷ if $m \geq \beta$ then return(m) // beta cut off, return m
 - end; return m
- **Algorithm $G2'(\text{position } p, \text{value } \alpha, \text{value } \beta)$ // min node**
 - determine the successor positions p_1, \dots, p_b
 - if $b = 0$, then return $f(p)$ else begin
 - ▷ $m := \beta$
 - ▷ for $i := 1$ to b do
 - ▷ $t := F2'(p_i, \alpha, m)$
 - ▷ if $t < m$ then $m := t$
 - ▷ if $m \leq \alpha$ then return(m) // alpha cut off, return m
 - end; return m

Alpha-beta pruning algorithm: Fail hard

Algorithm $F2(\text{position } p, \text{value } \alpha, \text{value } \beta, \text{integer } \text{depth})$

```

• determine the successor positions  $p_1, \dots, p_b$ 
• if  $b = 0$  // a terminal node
  or  $\text{depth} = 0$  // remaining depth to search
  or time is running up // from timing control
  or some other constraints are met // add knowledge here
• then return  $h(p)$  else
• begin
  ▷  $m := \alpha$ 
  ▷ for  $i := 1$  to  $b$  do
  ▷ begin
  ▷  $t := -F2(p_i, -\beta, -m, \text{depth} - 1)$ 
  ▷ if  $t > m$  then  $m := t$ 
  ▷ if  $m \geq \beta$  then return( $m$ ) // cut off, return  $m$  that is  $\geq \beta$ 
  ▷ end
• end
• return  $m$ 

```

Fail soft version

Algorithm $F3(\text{position } p, \text{value } \alpha, \text{value } \beta, \text{integer } \text{depth})$

```

• determine the successor positions  $p_1, \dots, p_b$ 
• if  $b = 0$  // a terminal node
  or  $\text{depth} = 0$  // remaining depth to search
  or time is running up // from timing control
  or some other constraints are met // add knowledge here
• then return  $h(p)$  else
• begin
  ▷  $m := -\infty$  // soft initial value
  ▷ for  $i := 1$  to  $b$  do
  ▷ begin
  ▷  $t := -F3(p_i, -\beta, -\max\{m, \alpha\}, \text{depth} - 1)$ 
  ▷ if  $t > m$  then  $m := t$  // the returned value is "used"
  ▷ if  $m \geq \beta$  then return( $m$ ) // cut off
  ▷ end
• end
• return  $m$ 

```

Properties and comments

Properties:

- Assumptions: (1) $\alpha < \beta$ and (2) p is not a leaf.
- $F2(p, \alpha, \beta) = \alpha$ if $F(p) \leq \alpha$
- $F2(p, \alpha, \beta) = F(p)$ if $\alpha < F(p) < \beta$
- $F2(p, \alpha, \beta) \geq \beta$ and $F(p) \geq F2(p, \alpha, \beta)$ if $F(p) \geq \beta$
- $F2(p, -\infty, +\infty) = F(p)$

Comments:

- $F2(p, \alpha, \beta)$: find the best possible value according to a nega-max formula for the position p with the constraints that
 - If $F(p) \leq \alpha$, then $F2(p, \alpha, \beta)$ returns with the value α from a terminal position whose value is $\leq \alpha$.
 - If $F(p) \geq \beta$, then $F2(p, \alpha, \beta)$ returns a value $\geq \beta$ from a terminal position whose value is $\geq \beta$.
- An intermediate version.
 - The lower bound is hard, cannot be violated.
 - Easier to find the branch where the returned value is coming from.
 - Always return something better than expected, but never something worse!!
- For historical reason [Fishburn 1983][Knuth & Moore 1975], this is called fail hard.

Properties and comments

Properties:

- Assumptions (1) $\alpha < \beta$ and (2) p is not a leaf
- $F3(p, \alpha, \beta, \text{depth}) \leq \alpha$ and $F(p) \leq F3(p, \alpha, \beta, \text{depth})$ if $F(p) \leq \alpha$
- $F3(p, \alpha, \beta, \text{depth}) = F(p)$ if $\alpha < F(p) < \beta$
- $F3(p, \alpha, \beta, \text{depth}) \geq \beta$ and $F(p) \geq F3(p, \alpha, \beta, \text{depth})$ if $F(p) \geq \beta$
- $F3(p, -\infty, +\infty, \text{depth}) = F(p)$


$F3$ finds a "better" value when the value is out of the search window.

- Better means a tighter bound.
 - The bounds are soft, i.e., can be violated.
- When it is failed-high, $F3$ normally returns a value that is higher than that of $F1$ or $F2$.
 - Never higher than that of $F!$
- When it is failed-low, $F3$ normally returns a value that is lower than that of $F1$ or $F2$.
 - Never lower than that of $F!$


Example

Initial call: $F2'(\text{root}, -\infty, \infty)$

- $m = -\infty$
- call $G2'(\text{node } 1, -\infty, \infty)$
 - it is a terminal node
 - return value 15
- $t = 15$;
 - since $t > m$, m is now 15
- call $G2'(\text{node } 2, 15, \infty)$
 - call $F2'(\text{node } 2.1, 15, \infty)$
 - it is a terminal node; return 10
 - $t = 10$; since $t < \infty$, m is now 10
 - α is 15, m is 10, so we have an alpha cut off.
 - no need to call $F2'(\text{node } 2.2, 15, 10)$
 - return 10
 - ...



Fail soft version: Example



- Let the value of the leaf node W be u .
- If $u < \alpha$, then the returned value of A will be at least u .


Comparisons between $F2$ and $F3$

- Both versions find the corrected value v if v is within the window $[\alpha, \beta]$.
- Both versions scan the same set of nodes during searching.
 - If the returned value of a subtree is decided by a cut, then $F2$ and $F3$ return the same value.
- $F3$ provides more information when the true value is out of the pre-assigned search window.
 - Can provide a feeling on how bad or good the game tree is.
 - Use this "better" value to guide searching later on.
- $F3$ saves about 7% of time than that of $F2$ when a **transposition table** is used to save and re-use searched results [Fishburn 1983].
 - A transposition table is a data structure to record the results of previous searched results.
 - The entries of a transposition table can be efficiently accessed, i.e., read and write, during searching.
 - Need an efficient addressing scheme, e.g., hash, to translate between a position and its address.

Comments

- For historical reason, comparisons are made between $F2$ and $F3$, while we should compare $F1$ and $F3$.
 - To me, $F1$ fails really hard. $F2$ is only an intermediate version!
- What move ordering is good?
 - It may not be good to search the best possible move first.
 - It may be better to cut off a branch with more nodes first.
- How about the case when the tree is not uniform?
- What is the effect of using iterative-deepening alpha-beta cut off?
- How about the case for searching a game graph instead of a game tree?
 - Can some nodes be visited more than once?

$F2$ and $F3$: Example (1/2)




- Assume the node A can be reached from the starting position using path P_1 and path P_2 .
 - If W is visited first along P_1 with a bound of $[4000, 5000]$, and returns a value of 200, then
 - the returned value of W , 200, is stored into the transposition table.
 - If A is visited again along P_2 with a bound of $[390, 600]$, then a better value of previously stored value of W helps to decide whether the subtree rooted at W needs to be searched again.

References and further readings

- * D. E. Knuth and R. W. Moore. An analysis of alpha-beta pruning. *Artificial Intelligence*, 6:293–326, 1975.
- * John P. Fishburn. Another optimization of alpha-beta search. *SIGART Bull.*, (84):37–38, 1983.
- * J. Pearl. The solution for the branching factor of the alpha-beta pruning algorithm and its optimality. *Communications of ACM*, 25(8):559–564, 1982.

$F2$ and $F3$: Example (2/2)



- Fail soft version has a chance to record a better value to be used later when this position is revisited.
 - If A is visited again along P_2 with a bound of $[390, 600]$, then
 - it does not need to be searched again, since the previous stored value of W is -200 .
 - However, if the value of W is 450, then it needs to be searched again.
- The fail hard version does not store the returned value of W after its first visit since this value is less than α .